

Mohammed Shaheem Abdul Salam

LinkedIn: [linkedin.com/in/shaheem32](https://www.linkedin.com/in/shaheem32) | Email: ashaheem32@gmail.com

GitHub: github.com/ashaheem32 | Portfolio: shaheem.me | Mobile: +91 6282102601

SUMMARY

AI Engineer specializing in production retrieval-augmented generation (RAG) and LLM systems. Published IEEE researcher (2026) — built a citation-grounded legal assistant that raised citation accuracy from 42% to 87%, cutting hallucination by 79%. Shipped end-to-end AI products spanning vector-search pipelines (pgvector, ChromaDB), real-time computer-vision analytics (YOLOv8), and FastAPI backends, backed by 1,000+ LLM evaluation and RLHF tasks for frontier-model training.

EXPERIENCE

Outlier AI — Freelance AI/ML Engineer (LLM Evaluation & RLHF)

Aug 2024 – Dec 2025

- Completed 1,000+ LLM evaluation and RLHF tasks for frontier-model training pipelines — assessing and improving model outputs across code generation, reasoning, and prompt optimization, specializing in Python and backend engineering domains.
- Maintained a consistent 4–5 star quality rating across diverse AI and software engineering projects, demonstrating reliable judgment on model correctness, groundedness, and code quality at scale.

TECHNICAL SKILLS

Programming: Python, SQL, TypeScript, JavaScript

LLM Systems: RAG pipelines, LangChain, vector databases (pgvector, ChromaDB), embeddings, OpenAI/Claude APIs, agentic workflows, LLM evaluation & RLHF

ML & Computer Vision: PyTorch, Scikit-learn, OpenCV, YOLOv8, ByteTrack, MediaPipe, Pandas, NumPy

Backend & Infrastructure: FastAPI, RESTful APIs, Node.js, PostgreSQL, SQLite, Docker, AWS, Supabase, Pytest, Git/GitHub

Languages: English, Hindi, Malayalam, Tamil

RESEARCH PUBLICATION

JurisGPT: A Retrieval-Augmented Generation Based Legal Assistant for Indian Startups (Accepted, IEEE)

Jan 2026

- Built a RAG-based legal platform grounding LLM responses in 8+ Indian statutory codes (3,169 vector-indexed segments) using ChromaDB HNSW and GPT-4o-mini; achieved 87% citation accuracy vs. 42% baseline, reducing hallucination by 79%.
- Engineered a four-stage pipeline with sentence-boundary-aware chunking, 384-dim embeddings (all-MiniLM-L6-v2), and a three-tier fallback architecture ensuring service continuity with end-to-end latency of 3–7 seconds.

PROJECTS

Chatseeker — AI Chat Analysis Platform | [Link](#)

Jan 2026 – Present

Python, FastAPI, PostgreSQL (pgvector), Claude API

- Architected a multi-stage AI pipeline ingesting chat exports from 6 platforms (WhatsApp, Telegram, Instagram, Facebook, CSV, JSON) into a self-designed Universal Chat JSON schema, performing emotion classification, language detection, and automated translation of non-English messages.
- Built citation-grounded semantic search by embedding every message into a pgvector index, powering streaming natural-language Q&A where each answer cites the exact source messages retrieved.
- Delivered a seven-module analytics dashboard covering sentiment timelines with annotated peaks, TF-IDF vocabulary analysis, behavioral metrics (response-time gauges, conversation-initiation patterns), and emoji-sentiment correlation.

Store Intelligence — Real-Time Retail Analytics | [Link](#)

Mar 2026 – Apr 2026

Python, YOLOv8, ByteTrack, FastAPI, Docker

- Built a real-time retail analytics pipeline using YOLOv8 + ByteTrack to detect and track visitors from CCTV footage, emitting structured events via idempotent batch HTTP ingestion.
- Developed a FastAPI backend with SQL serving live metrics, conversion funnels, zone heatmaps, and a 4-detector anomaly engine, containerized with Docker Compose.
- Delivered a live Rich TUI dashboard with async polling and full integration test coverage using Pytest + ASGITransport.

CERTIFICATES

OCI Generative AI Professional — Oracle (Oct 2025) | [Link](#)

OCI Data Science Professional — Oracle (Jul 2025) | [Link](#)

EDUCATION

Lovely Professional University, Phagwara — B.Tech, CSE; CGPA: 7.00

Aug 2022 – May 2026